

STAT 431 — Applied Bayesian Analysis — Course Notes

# Introduction to Computational Methods Part 2

Spring 2019

# Gibbs Sampling Software

Software is available that automates the Gibbs sampling process and extends it to handle non-conjugacy.

Cowles illustrates use of **OpenBUGS**  
(<http://www.openbugs.net>).

We will instead use **JAGS**  
(<http://mcmc-jags.sourceforge.net/>).

JAGS and OpenBUGS are similar, but have some important differences.

## OpenBUGS:

- ▶ Point-and-click interface (Windows)
- ▶ Extensive online tutorial with many classic examples

## JAGS:

- ▶ Developed actively
- ▶ Cross-platform (easily)
- ▶ Web site more stable

Both can interface with R through appropriate packages.

JAGS has its own interface — a simple command line — but we will instead run it from within R using package `rjags`.

(Other JAGS-related R packages exist — see JAGS manual.)

To run JAGS (or OpenBUGS) you need:

- ▶ the data
- ▶ a specification of the model (including priors)
- ▶ the initial values (for all chains)

## R/JAGS Example 8.5:

Population Proportion

**BEWARE:** JAGS and OpenBUGS are not fully compatible — OpenBUGS code (in Cowles and on web sites) may need modifications to be used as JAGS code.

Download and use the current JAGS user manual to avoid problems.

# JAGS Model Specification

Every model is specified in a `model` block:

```
model {  
  ... statements ...  
}
```

Many statements have the form

```
variablename ~ distribution(params)
```



Example:

$$Y \mid \pi \sim \text{binomial}(n = 10, \pi)$$

$$\pi \sim \text{beta}(\alpha = 2, \beta = 1)$$

```
model {  
  y ~ dbin(pi, 10)  
  pi ~ dbeta(2, 1)  
}
```

- ▶ the order of the statements *does not* matter
- ▶ the order of the parameters *does* matter
- ▶ variables do not need to be “declared”

There is a **graph** for this model:

[ Draw model graph ... ]

The graph says that:

- ▶  $\pi$  and  $y$  are both specified as random variables
- ▶  $\pi$  is specified marginally (unconditionally)
- ▶  $y$  is specified conditionally on  $\pi$

Variables may be specified together as a **vector** or as an **array**.

Eg: If  $y$  is a vector,  $y[i]$  is its  $i$ th element.

Vectors can be handled easily using a for statement, e.g. if  $y$  has  $n$  elements,

```
for (i in 1:n) {  
    y[i] ~ distribution(params)  
}
```

Important:  $n$  must be a constant (non-random).

Example:

$$Y_1, \dots, Y_n \mid \mu, \tau^2 \sim \text{i.i.d. } N(\mu, 1/\tau^2)$$

$$\mu \sim N(0, 100)$$

$$\tau^2 \sim \text{gamma}(0.01, 0.01)$$

} independent

```
model {  
  for (i in 1:n) {  
    y[i] ~ dnorm(mu, tausq)  
  }  
  mu ~ dnorm(0, 0.01)  
  tausq ~ dgamma(0.01, 0.01)  
}
```

The model graph:

[ Draw model graph ... ]

A **plate** indicates that the variables on it are repeated as vector/array elements.

The length of a *data* vector *y* can be referenced as `length(y)`:

```
model {  
  for (i in 1:length(y)) {  
    y[i] ~ dnorm(mu, tausq)  
  }  
  ...  
}
```

This avoids the need to specify the length as a separate data value.

Improper priors, such as

$$\mu \sim 1 d\mu$$

are **NOT** allowed in JAGS.

You can instead use a “vague” proper prior, such as

$$\mu \sim \text{dnorm}(0, 0.000001)$$

though whether this is “vague” enough depends on the context.

What if we want inference about  $\sigma^2$  (not just  $\tau^2$ )?

We could use

```
model {  
  for (i in 1:length(y)) {  
    y[i] ~ dnorm(mu, tausq)  
  }  
  mu ~ dnorm(0, 0.000001)  
  tausq ~ dgamma(0.01, 0.01)  
  
  sigmasq <- 1/tausq  
}
```



A statement

```
variablename <- expression
```

represents a **deterministic** relationship: the variable is computed according to the expression.

The expression may use the four usual arithmetic symbols, the unary minus, the power symbol  $\wedge$ , and certain scalar functions — see JAGS documentation.

Deterministic relationships are *hollow* lines in the model graph:

[ Draw model graph ... ]

Important: Variables having data values are *not allowed* on the left-hand side of a deterministic relationship (`<-`).

Eg: Suppose data `y[i]` needs to be log-transformed to normality.

**NOT** allowed:

```
y[i] <- exp(z[i])  
z[i] ~ dnorm(mu, tausq)
```

Also **NOT** allowed (in the `model` block):

```
z[i] <- log(y[i])  
z[i] ~ dnorm(mu, tausq)
```

Use a preceding data block to transform data:

```
data {  
  for (i in 1:length(y)) {  
    z[i] <- log(y[i])  
  }  
}  
  
model {  
  for (i in 1:length(z)) {  
    z[i] ~ dnorm(mu, tausq)  
  }  
  ...  
}
```

The data block is executed (once) before the model is run.

## Example: Wikipedia Article Modifications

For  $n = 10$  random English Wikipedia articles, time (days) since last modification is recorded:

$$Y_1, \dots, Y_{10} \sim \text{i.i.d. (given parameters)}$$

We want to know about the *median* time, and perhaps the variation.

Preliminary analysis (Q-Q plot) suggests approximate normality on the *log* scale:

$$\ln(Y_1), \dots, \ln(Y_{10}) \mid \mu, \sigma^2 \sim \text{i.i.d. } N(\mu, \sigma^2)$$

Consider prior

$$\left. \begin{array}{l} \mu \sim \text{N}(0, \sigma_0^2) \\ \sigma^2 \sim \text{IG}(\alpha, \beta) \end{array} \right\} \text{independent}$$

We will try

$$\sigma_0^2 = 10000 \quad \alpha = 0.0001 \quad \beta = 0.0001$$

because they produce a “vague” prior.

We might want inference for  $\mu$ ,  $\sigma^2$ , and

$$e^{\mu} = (\text{conditional}) \text{ median of } Y$$

(Why is that the median?)

Remember, JAGS uses

$$\tau^2 = 1/\sigma^2$$

Graph:

[ Draw model graph ... ]

Note: Use a square (instead of a circle) for values that are “constants,” i.e. not on the left-hand side of  $\sim$  or  $<-$ .



## R/JAGS Example 8.6:

(Transformed) Normal Sample



# Convergence Assessment

- ▶ Initial Values

... should be **overdispersed** — some should be chosen far away from the values you expect to see for the parameters

► Trace (History) Plots

check for convergence time and speed of “mixing”

Eg: showing transient (initialization) effects

[ Draw trace plot ... ]

Eg: showing lack of convergence

[ Draw trace plot ... ]

Eg: showing slow mixing (which gives high MC error)

[ Draw trace plot ... ]

► Autocorrelation Plots

check for serial dependence

Eg: showing high correlations (which gives slow mixing)

[ Draw autocorrelation plot ... ]

## ► Gelman-Rubin Statistic and Plots

The **Gelman-Rubin statistic**, or “potential scale reduction factor,” helps to monitor convergence when using multiple chains with overdispersed starting points.

For a specific monitored quantity, it is essentially

$$R = \sqrt{\frac{\hat{V}}{W}}$$

where

$\hat{V}$  = an estimated variance of the quantity  
based on all chains (“pooled”)

$W$  = average of the estimated variances  
from each chain individually

Idea:

[ Draw trace plots ... ]

Recommend  $R < 1.05$  to declare convergence.



The coda function `gelman.plot` plots  $R$  versus iterations used, which makes it easier to read how many iterations should be burned.

[ Draw example plot ... ]

# Posterior Predictive Distributions

JAGS offers two ways to sample from a posterior predictive distribution:

- ▶ add a new variable to your model (e.g. “ynew”)
- ▶ add a “missing” value in your data (using “NA”)

## Example: Wikipedia Article Modifications (continued)

Want:

- ▶ a range containing the days since modification for 95% of all articles
- ▶ the estimated percentage of articles modified in the last 30 days

We will add a new variable called `ynew` and another variable indicating whether `ynew` is less than or equal to 30.

## R/JAGS Example 8.7:

(Transformed) Normal Sample  
with Posterior Prediction

Notice: Certain transformations (including `log`) are allowed on the left hand side of a deterministic relationship:

```
log(ynew) <- znew
```

Q: Why didn't we use the "missing" value method?

A: Because the transformation of `y[i]` in the data block (on the right hand side of a `<-` statement) requires all non-missing values.